

Data Structures


Lecture 8

Lazy Binomial Heaps

Fibonacci Heaps

Shiri Chechik, Or Zamir
Winter semester 2025-6


Binary Heap



Binary Heaps	
Insert	$O(\log n)$
Find-min	$O(1)$
Delete-min	$O(\log n)$
Decrease-key	$O(\log n)$
Meld / Join	$O(n)$

Heaps


Binomial Heap



Binomial Heaps	
Insert	$O(\log n)$
Find-min	$O(1)$
Delete-min	$O(\log n)$
Decrease-key	$O(\log n)$
Meld / Join	$O(\log n)$

Q


Binomial Heap Representation



“Structure V”
[Brown (1978)]

Q

Alternative Representation



Linking binomial trees

Function link(x, y)

```
if  $x.key > y.key$  then
     $x \leftrightarrow y$ 
 $y.next \leftarrow x.child$ 
 $x.child \leftarrow y$ 
return  $x$ 
```

Linking in first
representation

Function link(x, y)

```
if  $x.key > y.key$  then
     $x \leftrightarrow y$ 
if  $x.child = null$  then
     $y.next \leftarrow y$ 
else
     $y.next \leftarrow x.child.next$ 
     $x.child.next \leftarrow y$ 
 $x.child \leftarrow y$ 
return  $x$ 
```

Linking in second
representation


Lazy Binomial Heaps



Binomial Heaps

	Binary Heaps	Binomial Heaps	Lazy Binomial Heaps	Fibonacci Heaps
Insert	$O(\log n)$	\leftarrow	$O(1)$	
Find-min	$O(1)$	\leftarrow	\leftarrow	
Delete-min	$O(\log n)$	\leftarrow	\leftarrow	
Decrease-key	$O(\log n)$	\leftarrow	\leftarrow	
Meld / Join	$O(n)$	$O(\log n)$	$O(1)$	

Worst case Amortized



Intuition

Intuition in a nutshell:

Be less rigid:

- **Binomial heap:**
eagerly link heaps at meld/insert/delete-min
- **Lazy binomial heap:**
lazily defer linking until next **delete-min**

Laziness will turn out beneficial
(amortized).

Benefits of laziness

Lazy Insert

Add the new item to the list of roots (as B_0)
Update the pointer to root with minimal key

$O(1)$ worst case time

Lazy Meld

Concatenate the two lists of Binomial trees
Update the pointer to root with minimal key


$O(1)$ worst case time

Lazy Binomial Heap

- Binomial heap:
A list of heap ordered binomial trees,
at most one of each degree
(at most $O(\log n)$ trees)
- **Lazy** binomial heap:
A list of heap ordered binomial trees
~~at most one of each degree~~
(possibly even n trees of size 1)


Lazy Binomial Heap

An **arbitrary** list of heap-ordered binomial trees
+ pointer to root with minimal key




Lazy Binomial Heap - Intuition

- A lazy binomial heap can be wide but **not deep**



Worst Case for Delete-min

- Remove the minimum root and meld exposed trees to the heap in $O(1)$





- But ?? time to find the new minimum!



Amortized Delete-min

- Delete-Min is a good opportunity to restore order by linking trees of same degree
 - This is called consolidating / successive linking
 - We'll now show this makes Delete-Min run in $O(\log n)$ amortized cost.



Consolidating / Successive Linking




Consolidating / Successive Linking




Consolidating / Successive Linking




Consolidating / Successive Linking




Consolidating / Successive Linking




Consolidating / Successive Linking




Consolidating / Successive Linking



Consolidating / Successive Linking




Consolidating / Successive Linking



Consolidating / Successive Linking

At the end of the process, we obtain a **non-lazy** binomial heap


(at most $\lfloor \log_2 n \rfloor + 1$ trees, at most 1 of each degree)



Worst Case Complexity of Consolidating

- Worst case (process all trees + all linkings):

$$T_0 - 1 + \log_2 n + L \leq 2(T_0 + \log_2 n) = O(n)$$



$L \leq T_0 + \log n$
why?

Amortized Complexity of Consolidating

- So actual worst case cost:
(scaled) $T_0 + \log_2 n$
- Claim: amortized complexity $O(\log n)$
 - Intuition?
 - Who can “pay for” consolidation?
 - Potential function?

Amortized Complexity of Consolidating

$\Phi = \text{Number of Trees}$

$$\Delta\Phi = T_1 - T_0$$

#trees after the process = $O(\log n)$

- Amortized cost = **actual cost** + $\Delta\Phi$

$$= (T_0 + \log n) + (T_1 - T_0) \quad * \text{up to scaling}$$

$$= \log n + T_1$$

$$= O(\log n)$$

Lazy Binomial Heaps

- Inserts pays for Del-min:

	Actual cost	Potential: Δ Trees*	Amortized cost
Lazy Insert	$O(1)$	+1	$O(1)$
Find-min	$O(1)$	0	$O(1)$
Delete-min	$T_0 + \log n$	$T_1 - T_0$	$O(\log n)$
Decrease-key	$O(\log n)$	0	$O(\log n)$
Lazy Meld	$O(1)$	0	$O(1)$

* up to scaling

Lazy Binomial Heaps

	Binary Heaps	Binomial Heaps	Lazy Binomial Heaps	Fibonacci Heaps
Insert	$O(\log n)$	\leftarrow	$O(1)$	
Find-min	$O(1)$	\leftarrow	\leftarrow	
Delete-min	$O(\log n)$	\leftarrow	\leftarrow	
Decrease-key	$O(\log n)$	\leftarrow	\leftarrow	
Meld / Join	$O(n)$	$O(\log n)$	$O(1)$	



Worst case

Amortized




Fibonacci Heaps

[Fredman-Tarjan (1987)]

Fibonacci Heaps


	Binary Heaps	Binomial Heaps	Lazy Binomial Heaps	Fibonacci Heaps
Insert	$O(\log n)$	\leftarrow	$O(1)$	\leftarrow
Find-min	$O(1)$	\leftarrow	\leftarrow	\leftarrow
Delete-min	$O(\log n)$	\leftarrow	\leftarrow	\leftarrow
Decrease-key	$O(\log n)$	\leftarrow	\leftarrow	$O(1)$
Meld / Join	$O(n)$	$O(\log n)$	$O(1)$	\leftarrow




Intuition in a nutshell

- Decrease-key: we do not want to fix heap order **all the way up ($O(\log n)$)**.


Decrease-key in $O(1)$ time?



Decrease-key in $O(1)$ time?



Decrease-key in $O(1)$ time?




Decrease-key($x, Q, \Delta = 10$)

No heap order violation

Decrease-key($x, Q, \Delta = 20$)


Decrease-key in $O(1)$ time?



Heap order violation


Can we avoid the $O(\log n)$?

Decrease-key in O(1) time?




A crazy idea: **cut** the edge
(add x to root list)

Decrease-key in O(1) time?




A crazy idea: **cut** the edge
(add x to root list)

Decrease-key in O(1) time?




Involved trees no longer
binomial

Decrease-key in O(1) time?




Decrease-key($x, Q, \Delta = 10$)

Decrease-key in O(1) time?



Cut the edge


Decrease-key in O(1) time?



Cut the edge

Fibonacci Heaps


A list of heap-ordered ~~binomial~~ trees
+ pointer to root with minimal key



All operations, except decrease-key are
the same as in *lazy binomial heaps*


Note on Linking in Fibonacci Heaps

- Linking 2 trees (after Delete-min) is done the same as in lazy binomial heaps:
 - Link 2 trees of same degree
- Only difference: trees not necessarily binomial



Fibonacci Heap - Intuition


- In a Fibonacci heap we may get almost any tree



The Problem with Cuts: Wide Shallow Trees

Intuition in a Nutshell


With many cuts, we may get **shallow wide** trees:



Which operation suffers from this?

Intuition in a Nutshell (2)


With many cuts, we may also get **deep narrow trees**:



Currently this is OK since Decrease-key cuts edges


Simple cuts create wide shallow trees

- Recall: a binomial tree of degree k contains 2^k nodes, so all degrees = $O(\log n)$
- However, with cuts, we may get trees of degree k containing as few as $k+1$ nodes, so degrees may be $O(n)$!
- Previous analysis of Del-min breaks down



How to eliminate wide shallow trees

- We don't want a node of degree k to have only $O(k)$ descendants.



- We want it to have $\Omega(c^k)$ descendants, for some constant $c > 1$, so $k = O(\log_c n)$.

Eliminating Wide Shallow Trees: Cascading Cuts (via Decrease-Key)

How to eliminate wide shallow nodes

In Decrease-Key:

- When a node **loses 2nd child** cut it too and add to root list
 - so a non-root node cannot lose many children without becoming a root itself
- Then we can prove: a node of degree k in a Fibonacci Heap has at least ϕ^k nodes, so $k = O(\log n)$ again!


Cascading cuts (eliminate wide shallow nodes)

Desired property:


Each node, except root, loses at most one child

- Rule 1: lose 1 child and you are marked LOSER
- Rule 2: lose 2nd child and you're dumped into root list (and unmarked)


Cascading cuts




Cascading cuts




Cascading cuts




Cascading cuts




Cascading cuts



Cascading cuts




- cut of decrease key
- cut of losing 2nd child




marked node

Cascading cuts




- cut of decrease key
- cut of losing 2nd child
- marked node

Cascading cuts



- cut of decrease key
- cut of losing 2nd child
- marked node


Cascading cuts




- cut of decrease key
- cut of losing 2nd child
- marked node

Plan for the Rest of this Lecture

- 1) Cascading cuts indeed prevent wide shallow trees




- 2) Decrease-Key may trigger a cascade of $O(n)$ cuts, but only $O(1)$ amortized



1) Cascading Cuts
Prevent Wide Shallow Trees


Maximally “damaged” trees

Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible




Maximally “damaged” trees

Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible




Maximally “damaged” trees

Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible




Maximally “damaged” trees

Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible




Maximally “damaged” trees

Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible




Maximally “damaged” trees

Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible




Maximally “damaged” trees

Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible




Maximally “damaged” trees

Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible




Maximally “damaged” trees

Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible




Maximally “damaged” trees

Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible




Maximally “damaged” trees

Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible




Maximally “damaged” trees

Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible




Maximally “damaged” trees

Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible



Maximally “damaged” trees


Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible



— cut of decrease key

Maximally “damaged” trees

Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible




— cut or decrease key

80

Maximally “damaged” trees


Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible



— cut of decrease key


Maximally “damaged” trees

Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible




Maximally “damaged” trees

Let's take a binomial tree of degree $k = 5$ and make it lose as many descendants as possible




Maximally “damaged” trees

Note the recursive structure




Maximally “damaged” trees

Note the recursive structure




Maximally “damaged” trees

What if we linked another tree
(of same degree $k = 5$),
and cut as much as possible?



Degrees in Maximally “damaged” trees

Lemma 1: Let x be a node of degree k and let y_1, y_2, \dots, y_k be the current children of x , in the order in which they were linked to x . Then, the degree of y_i is at least $i-2$.



Degrees in Maximally “damaged” trees

Lemma 1: Let x be a node of degree k and let y_1, y_2, \dots, y_k be the current children of x , in the order in which they were linked to x . Then, the degree of y_i is at least $i-2$.

Proof: When y_i was linked to x , y_1, \dots, y_{i-1} were already children of x . At that time, the degree of x was $i-1$. This was also the degree of y_i (why?) As y_i is still a child of x , it lost at most one child.

Size of Maximally “damaged” trees


Lemma 2: A node of degree k in a Fibonacci Heap has at least $F_{k+2} \geq \phi^k$ descendants, including itself.

$$\begin{aligned} F_0 &= 0 & F_1 &= 1 \\ F_k &= F_{k-1} + F_{k-2}, \quad k \geq 2 & \phi &= \frac{1+\sqrt{5}}{2} \simeq 1.618 \end{aligned}$$

n	0	1	2	3	4	5	6	7	8	9
F_n	0	1	1	2	3	5	8	13	21	34

Size of Maximally “damaged” trees

Lemma 2: A node of degree k in a Fibonacci Heap has at least $F_{k+2} \geq \phi^k$ descendants, including itself.




$$F_k$$

+

$$F_{k+1}$$

$$= F_{k+2}$$

Induction



Size of Maximally “damaged” trees


Lemma 2: A node of degree k in a Fibonacci Heap has at least $F_{k+2} \geq \phi^k$ descendants, including itself.

Corollary: In a Fibonacci heap containing n items, all degrees are at most $\log_\phi n \leq 1.4404 \log_2 n$

Degrees are again $O(\log n)$

Fibonacci Heap - Intuition

- In a Fibonacci heap a tree cannot be too wide!



2) Decrease-key is $O(1)$
amortized

Amortized Cost of Decrease-key

- A decrease-key operation may trigger $O(n)$ cuts.
- However, we want to prove the cost of Decrease-key is $O(1)$ amortized
 - Who can pay for all these cuts?
 - How should Φ be defined?

Potential – first try


$$\Phi = \#\text{trees} + \#\text{marked nodes}$$

We need this to pay
for Delete-Min,
remember?


So Decrease-key
will pay for itself

- We currently have 2 sources to potential change. We have to see if previous amortized analysis has changed

Amortized Cost of Decrease-key



Amortized Cost of Decrease-key



- c new trees
- ≤ 1 new mark (last cut), and
- c or $c - 1$ marks removed


$$\underline{\Delta\Phi}$$

$$+c$$

$$\leq 2 - c$$

Amortized Cost of Decrease-key

	Actual cost	potential: Δ Trees	Potential: Δ Marks	Amortized cost
Decrease-key	c	$+c$	$\leq 2-c$	$O(c) = O(n)$



number of cuts
performed

Potential - Solution

Potential = #trees + 2·#marked nodes

note this 2

- We currently have 2 sources to potential change. We have to see if previous amortized analysis has changed

Amortized Cost of Decrease-key

	Actual cost	potential: Δ Trees	Potential: $2 \cdot \Delta$ Marks	Amortized cost
Decrease-key	c	$+c$	$\leq 2 \cdot (2 - c)$	$O(1)$



number of cuts
performed




Fibonacci heaps operations

- Inserts pays for Del-min, Decrease-key pays for itself:

	Actual cost	Potential: $\Delta \text{ Trees}^*$	Potential: $2 \cdot \Delta \text{ Marks}^*$	Amortized cost
Insert	$O(1)$	+1	0	$O(1)$
Find-min	$O(1)$	0	0	$O(1)$
Delete-min	$T_0 + \log n$	$T_1 - T_0$	≤ 0	$O(\log n)$
Decrease-key	c	$+c$	$\leq 2(2-c)$	$O(1)$
Meld	$O(1)$	0	0	$O(1)$


* up to scaling

Putting it all together (account method)




$$\text{Potential} = \#\text{trees} + 2 \, \#\text{marked nodes} = 1 + 2 \cdot 0 = 1$$

Putting it all together (account method)



$$\text{Potential} = \#\text{trees} + 2 \, \#\text{marked nodes} = 2 + 2 \cdot 1 = 4$$


Putting it all together (account method)



Decrease-key($x, Q, \Delta = 10$)

$$\text{Potential} = \#\text{trees} + 2 \cdot \#\text{marked nodes} = 2 + 2 \cdot 1 = 4$$


Putting it all together (account method)



Decrease-key($x, Q, \Delta = 10$)


Potential = #trees + 2 #marked nodes = $3 + 2 \cdot 1 = 5$

Putting it all together (account method)



$$\text{Potential} = \#\text{trees} + 2 \cdot \#\text{marked nodes} = 4 + 0 \cdot 1 = 4$$

Putting it all together (account method)



$$\text{Potential} = \#\text{trees} + 2 \cdot \#\text{marked nodes} = 4 + 0 \cdot 1 = 4$$



Heaps / Priority queues

	Binary Heaps	Binomial Heaps	Lazy Binomial Heaps	Fibonacci Heaps
Insert	$O(\log n)$	←	$O(1)$	←
Find-min	$O(1)$	←	←	←
Delete-min	$O(\log n)$	←	←	←
Decrease-key	$O(\log n)$	←	←	$O(1)$
Meld / Join	$O(n)$	$O(\log n)$	$O(1)$	←

[]

Worst case

Amortized



Summary: Life is all about Tradeoffs

- If we impose no structural constraints on our trees, then trees of large degree may have too few nodes. This leads to wrecking our runtime bounds for extract-min.
- If we impose too many structural constraints on our trees, then we have to spend too much time fixing up trees. This leads to decrease-key taking too long.
- Fibonacci heaps strike a balance
 - If we do a few decrease-keys, then the tree won't lose “too many” nodes.
 - If we do many decrease-keys, the information slowly propagates to the root (its degree slowly decreases).

Fibonacci Heaps: theory vs. practice

- Theoretically they look good
- Practically:
 - Less efficient than “simpler” heaps
 - Complicated to code and use a lot of memory for each node (see next slide)
 - Still $O(n)$ worst case for Delete-min

Q


Fibonacci heap representation

min

Explicit degrees


Doubly linked lists at each level

Parent points to an arbitrary child



4 pointers + degree + mark bit per node

Fibonacci heap representation



4 pointers + degree + mark bit per node

Cascading cuts

Function $\text{cut}(x, y)$

```
x.parent ← null  
x.mark ← 0  
y.rank ← y.rank - 1  
if x.next = x then  
| y.child ← null  
else  
| y.child ← x.next  
| x.prev.next ← x.next  
| x.next.prev ← x.prev
```

Cut x from its parent y

Function $\text{cascading-cut}(x, y)$

```
cut(x, y)  
if y.parent ≠ null then  
| if y.mark = 0 then  
| | y.mark ← 1  
| else  
| | cascading-cut(y, y.parent)
```

Perform a cascading-cut
process starting at x

Consolidating / Successive linking

Function consolidate(x)
to-buckets(x)
return from-buckets()

Function to-buckets(x)

for $i \leftarrow 0$ to $\log_\phi n$ **do**
 $B[i] \leftarrow \text{null}$

$x.\text{prev}.\text{next} \leftarrow \text{null}$

while $x \neq \text{null}$ **do**

$y \leftarrow x$

$x \leftarrow x.\text{next}$

while $B[y.\text{rank}] \neq \text{null}$ **do**

$y \leftarrow \text{link}(y, B[y.\text{rank}])$

$B[y.\text{rank} - 1] \geq \text{null}$

$B[y.\text{rank}] \leftarrow y$

Function from-buckets()

$x \leftarrow \text{null}$

for $i \leftarrow 0$ to $\log_\phi n$ **do**
 if $B[i] \neq \text{null}$ **then**
 if $x = \text{null}$ **then**
 $x \leftarrow B[i]$
 $x.\text{next} \leftarrow x$
 $x.\text{prev} \leftarrow x$
 else
 insert-after($x, B[i]$)
 if $B[i].\text{key} < x.\text{key}$ **then**
 $x \leftarrow B[i]$

return x

Heaps: famous last words...

- Binary heaps, binomial heaps, and Fibonacci heaps are all **inefficient** in their support of **Search**
 - Operations such as Decrease-key Delete-min require a **pointer** to the node
- Min vs. max
- A highly recommended summary:
https://en.wikipedia.org/wiki/Fibonacci_heap